

# Data 301 Data Analytics

## Database Part II: SQL

**Dr. Irene Vrbik**

University of British Columbia Okanagan  
irene.vrbik@ubc.ca

# Intro

- ▶ So far we've learnt how to manipulate data stored in a relational database using:

**CREATE** *to create new tables*

**INSERT** *to insert new rows*

**UPDATE** *to edit rows in a table*

**ALTER** *to edit columns in a table*

**DELETE** *for deleting rows*

**DROP TABLE** *for deleting tables*

- ▶ Today we'll look at some examples for *querying* data from a database.

## SQL

# SQL Queries using SELECT

A query in SQL has the form:

```
SELECT <list of columns separated by commas>
FROM <list of tables>
WHERE <filter conditions>
GROUP BY <columns>
ORDER BY <columns> ASC/DESC
```

Notes:

1. Separate the list of columns/expressions and list of tables by commas.
2. The "\*" is used to select all columns.
3. Only SELECT required. FROM, WHERE, GROUP BY, ORDER BY are optional.

## Doing a query in SQL mode

- ▶ To see how to do a basic **SELECT** query in Microsoft Access, see 1:35-2:25 in [this](#) YouTube video.
- ▶ To see how to do a basic **SELECT** query in LibreOffice base, see [this](#) YouTube demo.

# Example Data

## Example Data

emp Table

eno	ename	bdate	title	salary	supereno	dno
E1	J. Doe	01-05-75	EE	30000	E2	null
E2	M. Smith	06-04-66	SA	50000	E5	D3
E3	A. Lee	07-05-66	ME	40000	E7	D2
E4	J. Miller	09-01-50	PR	20000	E6	D3
E5	B. Casey	12-25-71	SA	50000	E8	D3
E6	L. Chu	11-30-65	EE	30000	E7	D2
E7	R. Davis	09-08-77	ME	40000	E8	D1
E8	J. Jones	10-11-72	SA	50000	null	D1

proj Table

pno	pname	budget	dno
P1	Instruments	150000	D1
P2	DB Develop	135000	D2
P3	Budget	250000	D3
P4	Maintenance	310000	D2
P5	CAD/CAM	500000	D2

workson Table

eno	pno	resp	hours
E1	P1	Manager	12
E2	P1	Analyst	24
E2	P2	Analyst	6
E3	P3	Consultant	10
E3	P4	Engineer	48
E4	P2	Programmer	18
E5	P2	Manager	24
E6	P4	Manager	48
E7	P3	Engineer	36

dept Table

dno	dname	mgreno
D1	Management	E8
D2	Consulting	E7
D3	Accounting	E5
D4	Development	null

## SQL: Retrieving Only Some Columns

The **projection operation** creates a new table that has some of the columns of the input table.

In SQL, provide the table in the **FROM** clause and the fields in the output in the **SELECT**.

Example: Return only the eno field from the emp table:

```
SELECT eno FROM emp
```

emp Table

eno	ename	bdate	title	salary	supereno	dno
E1	J. Doe	01-05-75	EE	30000	E2	null
E2	M. Smith	06-04-66	SA	50000	E5	D3
E3	A. Lee	07-05-66	ME	40000	E7	D2
E4	J. Miller	09-01-50	PR	20000	E6	D3
E5	B. Casey	12-25-71	SA	50000	E8	D3
E6	L. Chu	11-30-65	EE	30000	E7	D2
E7	R. Davis	09-08-77	ME	40000	E8	D1
E8	J. Jones	10-11-72	SA	50000	null	D1



Result

eno
E1
E2
E3
E4
E5
E6
E7
E8

# SQL Projection Examples

emp Table

eno	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

```
SELECT eno, ename  
FROM emp
```

eno	ename
E1	J. Doe
E2	M. Smith
E3	A. Lee
E4	J. Miller
E5	B. Casey
E6	L. Chu
E7	R. Davis
E8	J. Jones

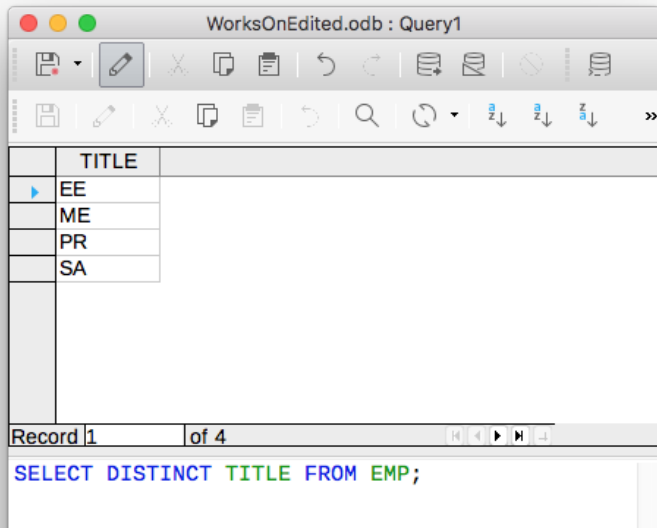
```
SELECT title  
FROM emp
```

title
EE
SA
ME
PR
SA
EE
ME
SA

## Notice

1. Duplicates are not removed during SQL projection.
2. **SELECT \*** will return all columns.

To return only the distinct values from the previous example, use **DISTINCT**:



WorksOnEdited.odb : Query1

	TITLE
▶	EE
	ME
	PR
	SA

Record 1 of 4

```
SELECT DISTINCT TITLE FROM EMP;
```



### Example 5.1

Given this table and the following query:

```
Select eno, ename, salary  
From emp
```

How many columns are returned?

- A) 0    B) 1    C) 2    D) 3    E) 4

emp Table

<u>eno</u>	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

## Answer

Given this table and the following query:

```
Select eno, ename, salary  
From emp
```

How many columns are returned?

- A) 0    B) 1    C) 2    D) 3    E) 4

emp Table

<u>eno</u>	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

### Example 5.2

Given this table and the following query:

```
Select salary  
From emp
```

How many *rows* are returned?

- A) 0    B) 1    C) 2    D) 4    E) 8

emp Table

<u>eno</u>	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

## Answer

Given this table and the following query:

```
Select salary  
From emp
```

How many *rows* are returned?

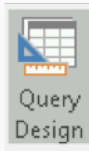
- A) 0    B) 1    C) 2    D) 4    E) 8

emp Table

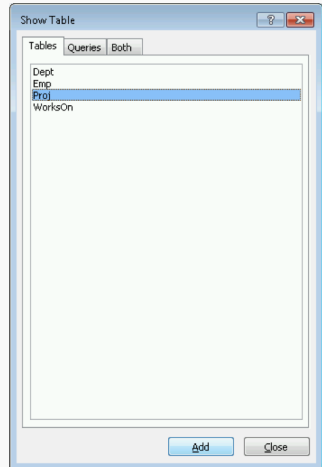
<u>eno</u>	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

# Building a **SELECT** SQL Query in Microsoft Access

Under **Create Tab**, click on **Query Design**.



Access will pop-up a window asking what table(s) you wish to query. Select one or more.



# Microsoft Access Query Interface

switch  
view  
button

The screenshot shows the Microsoft Access Query Design View for a query named 'SampleQuery'. The interface includes a ribbon with tabs: File, Home, Create, External Data, Database Tools, and Design. The Design tab is active, showing tools for inserting, deleting, and showing rows and columns, as well as a 'Show/Hide' button. The 'All Tables' pane on the left lists tables: Dept, Emp, Proj, and WorksOn. The main area displays three tables: Proj, WorksOn, and Emp, connected by lines representing relationships. Below the tables is a criteria table with columns for fields and rows for sorting and criteria.

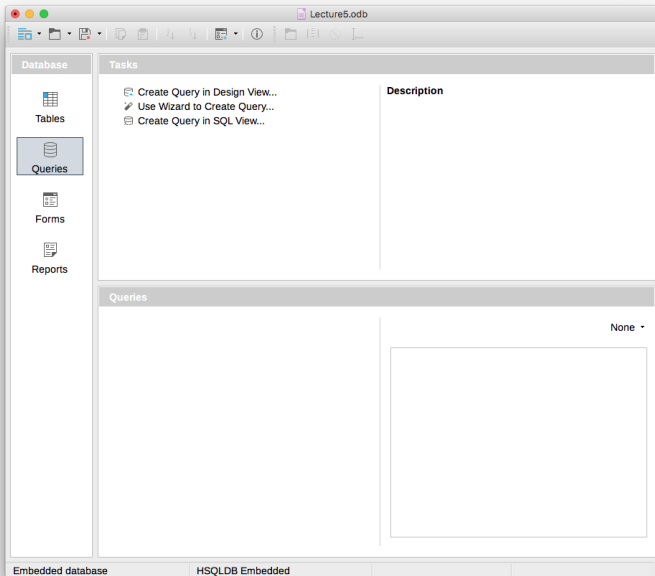
Tables are boxes. Relationships are lines.

fields in result sorting

selection criteria

Field:	ename	pname	hours	title
Table:	Emp	Proj	WorksOn	Emp
Sort:	Ascending			
Show:	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
Criteria:			>10	'EE' Or 'SA'
or:				

# LibreOffice Query Interface



# SQL

## LibreOffice Query Interface

Lecture5.odb : Query1

```
graph LR; proj --> WorksOn; WorksOn --> EMP;
```

The diagram shows three tables: **proj**, **WorksOn**, and **EMP**.


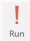
- proj** table fields: pno, pname, budget, dno.
- WorksOn** table fields: eno, pno, resp, hours.
- EMP** table fields: ENO, ENAME, BDATE, TITLE, SALARY, SUPERENO, DNO.

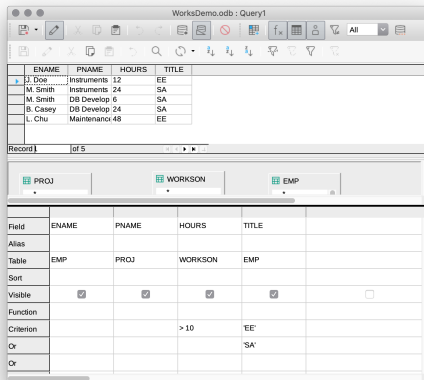
The relationships are as follows:

- proj** is connected to **WorksOn** via the **pno** field.
- WorksOn** is connected to **EMP** via the **eno** field.

Field	ENAME	pname	hours	TITLE					
Alias									
Table	EMP	proj	WorksOn	EMP					
Sort									
Visible	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
Function									
Criterion			> 10	'EE'					
Or				'SA'					
Or									
Or									
Or									



- ▶ If you hit the Run command (  in LibreOffice base,  in Access) and results will instantly appear in a data sheet view.
- ▶ We can save these results in Access by clicking the "Save" icon Toolbar. Give your query an identifying name, eg "Query1" for assignment question 1.



# Microsoft Access Data Sheet View

SampleQuery - Access

File Home Create External Data Database Tools Tell me what you want to do... Sign in

View Paste Filter Sort & Filter Records Find Window Text Formatting

**All Tables**

- Dept
  - Dept : Table
- Emp
  - Emp : Table
  - SampleQuery
- Proj
  - Proj : Table
  - SampleQuery
- WorksOn
  - WorksOn : Table
  - SampleQuery

ename	pname	hours	title
B. Casey	DB Develop	24 SA	
J. Doe	Instruments	12 EE	
L. Chu	Maintenance	48 EE	
Rich Guy	Instruments	24 SA	
*			

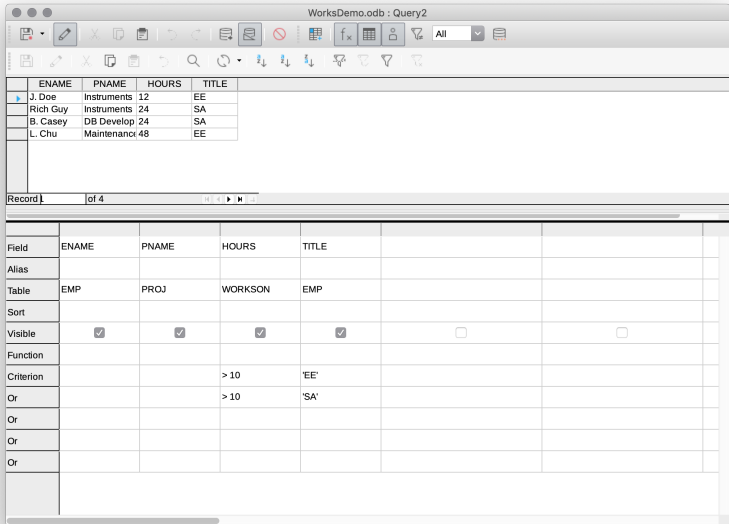
Record: 5 of 5 No Filter Search

Ready Num Lock SQL

# LibreOffice Preview

To view the table: **View** → **Preview** or **F5** or Run the Query

WorksDemo.odb : Query2



	ENAME	PNAME	HOURS	TITLE
▶	J. Doe	Instruments	12	EE
	Rich Guy	Instruments	24	SA
	B. Casey	DB Develop	24	SA
	L. Chu	Maintenanci	48	EE

Record 1 of 4

Field	ENAME	PNAME	HOURS	TITLE			
Alias							
Table	EMP	PROJ	WORKSON	EMP			
Sort							
Visible	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	
Function							
Criterion			> 10	'EE'			
Or			> 10	'SA'			
Or							
Or							
Or							

## SQL Query View in LibreOffice

The design view window in Microsoft Access and LibreOffice provide an easy way of creating SQL queries

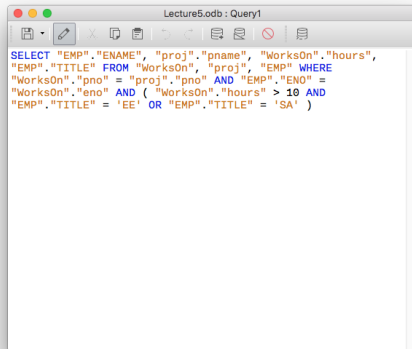
This may be referred to as QBE or Query By Example

Under the hood, these programs are creating SQL code and runs it to give us our result set.

To view that code in LibreOffice, we go to **View** → **Switch Design View On/Off**.

- ▶ 'Off' will show the query in SQL
- ▶ 'On' will show the query graphical

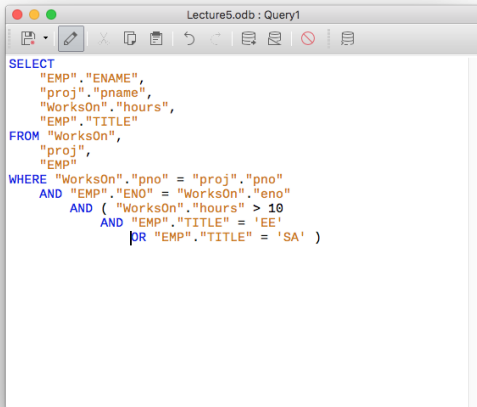
# LibreOffice Query Views

A screenshot of the LibreOffice Query Designer window. The title bar reads "Lecture5.odb : Query1". The window contains a SQL query with syntax highlighting. The query is: 

```
SELECT "EMP"."ENAME", "proj"."pname", "WorksOn"."hours",  
"EMP"."TITLE" FROM "WorksOn", "proj", "EMP" WHERE  
"WorksOn"."pno" = "proj"."pno" AND "EMP"."ENO" =  
"WorksOn"."eno" AND ( "WorksOn"."hours" > 10 AND  
"EMP"."TITLE" = 'EE' OR "EMP"."TITLE" = 'SA' )
```

Notice how the SQL code is not formatted very pretty. We can edit the white space, however, every time we close it, it will just get ugly again.

# LibreOffice Query Views

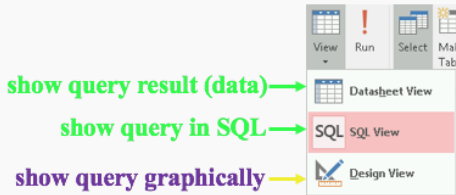
A screenshot of the LibreOffice Query Designer window. The title bar reads "Lecture5.odb : Query1". The window contains an SQL query with the following text:

```
SELECT
    "EMP"."ENAME",
    "proj"."pname",
    "WorksOn"."hours",
    "EMP"."TITLE"
FROM    "WorksOn",
        "proj",
        "EMP"
WHERE   "WorksOn"."pno" = "proj"."pno"
AND     "EMP"."ENO" = "WorksOn"."eno"
AND (   "WorksOn"."hours" > 10
        AND "EMP"."TITLE" = 'EE'
        OR "EMP"."TITLE" = 'SA' )
```

Notice how the SQL code is not formatted very pretty. We can edit the white space, however, every time we close it, it will just get ugly again.

# Microsoft Access Queries in SQL View

You may view your data, your query graphically, or your query in SQL.



For exam purposes, we will be needing to familiarize ourselves more with the SQL View.

# LibreOffice base Queries in SQL View

The screenshot shows the LibreOffice Base application interface. On the left, the 'Database' sidebar has 'Queries' selected. The 'Tasks' pane shows three options: 'Create Query in Design View...', 'Use Wizard to Create Query...', and 'Create Query in SQL View...'. An orange circle highlights the 'Queries' icon in the sidebar, with an arrow pointing to the 'Create Query in SQL View...' task. Another orange circle highlights the 'Create Query in SQL View...' task, with an arrow pointing to a floating window titled 'WorksOnEdited.odb : Query2'. This window displays a table with 5 records and 4 columns: PNO, PNAME, BUDGET, and DNO. An orange circle highlights the 'View Results' icon (a document with a magnifying glass) in the window's toolbar, with an arrow pointing to the table. Below the table, the SQL statement 'SELECT \* FROM PROJ;' is visible. An orange circle highlights the 'challenge1' query name in the 'Queries' list, with an arrow pointing to the text 'we can save the results of a query'.

clicking this brings up this window

click this icon to view the results of this query

we can save the results of a query

PNO	PNAME	BUDGET	DNO
P1	Instruments	150000.00	D1
P2	DB Develop	135000.00	D2
P3	Budget	250000.00	D3
P4	Maintenanci	310000.00	D2
P5	CAD/CAM	500000.00	D1

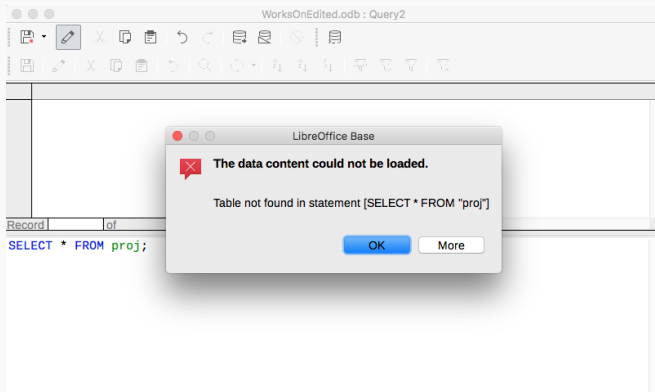
```
SELECT * FROM PROJ;
```



# SQL Query View in LibreOffice

## Warning

Unlike the SQL statements from last lecture, the SQL view mode for queries is case sensitive. That is, it will not convert all of our lower case to upper case text as it did before (annoying I know!)



## Try it: SQL **SELECT** and Projection

### Example 5.3

Using the proj table, write these three queries:

- ▶ Show all rows and all columns.
- ▶ Show all rows but only the pno column.
- ▶ Show all rows but only the pno and budget columns.

## Retrieving only some of the rows

The **selection operation** creates a new table with some of the rows of the input table.

A **condition** specifies which rows are in the new table. This condition is similar to a filter in Excel.

Eg. the following algorithm scans each tuple and checks if it satisfies the condition in the **WHERE** clause.

```
SELECT * FROM proj WHERE dno = 'D2';
```

**proj Table**

pno	pname	budget	dno
P1	Instruments	150000	D1
P2	DB Develop	135000	D2
P3	Budget	250000	D3
P4	Maintenance	310000	D2
P5	CAD/CAM	500000	D2

**Result**

pno	pname	budget	dno
P2	DB Develop	135000	D2
P4	Maintenance	310000	D2
P5	CAD/CAM	500000	D2



## Selection Conditions

- ▶ The condition in a selection statement specifies which rows are included.
- ▶ It has the general form of an if statement.
- ▶ The condition may consist of attributes, constants, comparison operators (<, >, =, !=, <=, >=), and logical operators (AND, OR, NOT).
- ▶ To check for NULLs, use IS NULL<sup>1</sup> which is different than checking if its equal to the empty string = ''. General syntax:

```
SELECT column_names  
FROM table_name  
WHERE column_name IS NULL;
```

---

<sup>1</sup>IS NOT NULL is used to check for non null values

# SQL Selection Examples

## emp Table

eno	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

```
SELECT *
FROM emp
WHERE title = 'EE'
```

eno	ename	title	salary
E1	J. Doe	EE	30000
E6	L. Chu	EE	30000

```
SELECT *
FROM emp
WHERE salary > 35000
OR title = 'PR'
```

eno	ename	title	salary
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

## Selection Question

### Example 5.4

Given the emp table and the following query, how many rows are returned?

```
SELECT *  
FROM emp  
WHERE title = 'SA'
```

emp Table

eno	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

- A) 0      B) 1      C) 2      D) 3      E) >3

# Selection Question

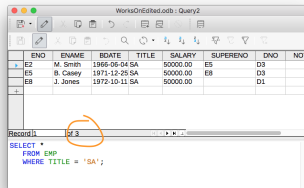
## Answer

Given the emp table and the following query, how many rows are returned?

```
SELECT *
FROM emp
WHERE title = 'SA'
```

emp Table

eno	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000



A) 0    B) 1    C) 2    D) 3    E) >3

## Selection Question2

### Example 5.5

Given the emp table and the following query, how many rows are returned?

```
SELECT *  
FROM emp  
WHERE salary > 40000  
       OR title='PR'
```

- A) 0
- B) 1
- C) 3
- D) 4
- E) > 4

emp Table

eno	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000



## Selection Question 2

### Answer

Given the emp table and the following query, how many rows are returned?

```
SELECT *
FROM emp
WHERE salary > 40000
      OR title='PR'
```

- A) 0
- B) 1
- C) 3
- D) 4
- E) > 4

### emp Table

eno	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

WorksOnEdited.odb : Query1

ENO	ENAME	BDATE	TITLE	SALARY	SUPERENO	DNO
E2	M. Smith	1966-06-04	SA	50000.00	E5	D3
E4	J. Miller	1950-09-01	PR	20000.00	E6	D3
E5	B. Casey	1971-12-25	SA	50000.00	E8	D3
E8	J. Jones	1972-10-11	SA	50000.00		D1

Record 1 of 4

```
SELECT *
FROM EMP
WHERE SALARY > 40000
      OR TITLE='PR';
```

## Selection Question 3

### Example 5.6

Given the emp table and the following query, how many rows are returned?

```
SELECT *  
FROM emp  
WHERE salary >= 40000  
      AND ename > 'C'
```

- A) 0
- B) 1
- C) 2
- D) 3
- E)  $\geq 4$

emp Table

eno	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

# Selection Question 3

## Answer

Given the emp table and the following query, how many rows are returned?

```
SELECT *
FROM emp
WHERE salary >= 40000
      AND ename > 'C'
```

- A) 0
- B) 1
- C) 2
- D) 3
- E)  $\geq 4$

emp Table

eno	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

The screenshot shows a database application window titled 'WorksOn/Edited.odb : Query1'. It displays a table with columns: ENO, ENAME, BOATE, TITLE, SALARY, SUPERENO, and DNO. The data rows are:

ENO	ENAME	BOATE	TITLE	SALARY	SUPERENO	DNO
E2	M. Smith	1968-06-04	SA	50000.00	E5	D3
E7	R. Davis	1977-09-08	ME	40000.00	E8	D1
E8	J. Jones	1972-10-11	SA	50000.00		D1

Below the table, it shows 'Record 1 of 3'. A query window is open with the following SQL:

```
SELECT *
FROM EMP
WHERE SALARY >= 40000
AND ENAME > 'C'
```

## Try It: SQL `SELECT` and Filtering Rows

### Example 5.7

Using the proj table, write these three queries:

- ▶ Return all projects with budget > \$250000.
- ▶ Show the pno and pname for projects in dno = 'D1'.
- ▶ Show pno and dno for projects in dno='D1' or dno='D2'.

## Join Example for Combining Tables

A **join** combines two tables by matching columns in each table.

eno	pno	resp	dur
E1	P1	Manager	12
E2	P1	Analyst	24
E2	P2	Analyst	6
E3	P4	Engineer	48
E5	P2	Manager	24
E6	P4	Manager	48
E7	P3	Engineer	36
E7	P4	Engineer	23

```
SELECT *
FROM WorksOn
INNER JOIN Proj
    ON WorksOn.pno = Proj.pno
```

**Figure:** workson

pno	pname	budget
P1	Instruments	150000
P2	DB Develop	135000
P3	CAD/CAM	250000
P4	Maintenance	310000
P5	CAD/CAM	500000

eno	pno	resp	dur	P.pno	pname	budget
E1	P1	Manager	12	P1	Instruments	150000
E2	P1	Analyst	24	P1	Instruments	150000
E2	P2	Analyst	6	P2	DB Develop	135000
E3	P4	Engineer	48	P4	Maintenance	310000
E5	P2	Manager	24	P2	DB Develop	135000
E6	P4	Manager	48	P4	Maintenance	310000
E7	P3	Engineer	36	P3	CAD/CAM	250000
E7	P4	Engineer	23	P4	Maintenance	310000

**Figure:** proj

The general syntax is:

```
SELECT <columns>
FROM
R
<type of join> JOIN
S
ON R.<colname>= S.<colname>;
```

Since joining tables often result in repeated field (ie. columns) we distinguish between them using `table_name.column_name`.

N.B. the related columns will not necessarily have the same name (but oftentimes they will!)

## Joining Tables

When connecting tables **R** and **S**, there are four types of joins:

**(INNER) JOIN** row in result for each row of R that matches a row of S

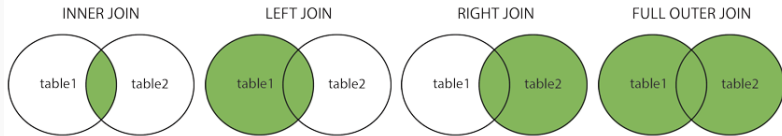
**LEFT (OUTER) JOIN** row in result for each row of R that matches a row of S OR a row of R that does not match anything in S

**RIGHT (OUTER) JOIN** row in result for each row of R that matches a row of S OR a row of S that does not match anything in R

**FULL OUTER JOIN** row in result for each row of R that matches a row of S OR a row of R that does not match anything in S OR a row of S that does not match anything in R

w3schools is another good resource for learning SQL. Here is a helpful representation of the different types of joins (where R = table1 and S=right table2).

- **(INNER) JOIN**: Returns records that have matching values in both tables
- **LEFT (OUTER) JOIN**: Returns all records from the left table, and the matched records from the right table
- **RIGHT (OUTER) JOIN**: Returns all records from the right table, and the matched records from the left table
- **FULL (OUTER) JOIN**: Returns all records when there is a match in either left or right table





# Join Example

**SELECT \* FROM Boys <> JOIN Girls ON Boys.Bid = Girls.Gid;**

Boys

Bid	BoyName
1	Joe
2	Steve
3	Fred
5	James

Boys INNER JOIN Girls

Bid	BoyName	Gid	GirlName
2	Steve	2	Jane
5	James	5	Fran

Boys LEFT OUTER JOIN Girls

Bid	BoyName	Gid	GirlName
1	Joe		
2	Steve	2	Jane
3	Fred		
5	James	5	Fran

Boys FULL OUTER JOIN Girls

Bid	BoyName	Gid	GirlName
1	Joe		
2	Steve	2	Jane
3	Fred		
		4	Sarah
5	James	5	Fran
		6	Julie

Girls

Gid	GirlName
2	Jane
4	Sarah
5	Fran
6	Julie

Boys RIGHT OUTER JOIN Girls

Bid	BoyName	Gid	GirlName
2	Steve	2	Jane
		4	Sarah
5	James	5	Fran
		6	Julie

## Join Query with Selection Example

You can use join, selection, and projection in the same query.  
Recall:

- ▶ **projection** returns *columns* listed in **SELECT**
- ▶ **selection** filters out *rows* using condition in **WHERE**, and
- ▶ **join** combines *tables* in **FROM** using a condition.

### Example 5.8

Return the employee names who are assigned to the 'Management' department.

## emp Table

eno	ename	bdate	title	salary	supereno	dno
E1	J. Doe	01-05-75	EE	30000	E2	null
E2	M. Smith	06-04-66	SA	50000	E5	D3
E3	A. Lee	07-05-66	ME	40000	E7	D2
E4	J. Miller	09-01-50	PR	20000	E6	D3
E5	B. Casey	12-25-71	SA	50000	E8	D3
E6	L. Chu	11-30-65	EE	30000	E7	D2
E7	R. Davis	09-08-77	ME	40000	E8	D1
E8	J. Jones	10-11-72	SA	50000	null	D1

## dept Table

<u>dno</u>	dname	mgreno
D1	Management	E8
D2	Consulting	E7
D3	Accounting	E5
D4	Development	null

**SELECT** ename  
**FROM** emp INNER JOIN dept  
           ON emp.dno = dept.dno  
**WHERE** dname = 'Management';

Projection: only name field in result  
 tables in query joined together  
 Selection: filter rows

## Result

ename
R. Davis
J. Jones

## Ordering Result Data

The query result returned is not ordered on any column by default. We can order the data using the **ORDER BY** clause:

```
SELECT ename, salary, bdate
FROM emp
WHERE salary > 30000
      ORDER BY salary DESC, ename ASC;
```

- ▶ **ASC** sorts the data in ascending order, and **DESC** sorts it in descending order. The default is **ASC**.
- ▶ The order of sorted attributes matters; namely, the first column specified is sorted on first, then the second column is used to break any ties, etc.

In LibreOffice the order is done Left to right. To ensure that salary gets ordered first, add another hidden (Deselect **Visible**) column with `ename` to the right.

DatabasesII.odb : Query2

	ENAME	SALARY	BDATE
	B. Casey	50000.00	1971-12-25
	J. Jones	50000.00	1972-10-11
	M. Smith	50000.00	1966-06-04
	A. Lee	40000.00	1966-07-05
	R. Davis	40000.00	1977-09-08

Record 1 of 5

EMP

Field	ENAME	SALARY	BDATE	ENAME			
Alias							
Table	EMP	EMP	EMP	EMP			
Sort		descending		ascending			
Visible	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
Function							
Criterion		>30000					
Or							
Or							

## LIMIT and OFFSET

If you only want the first  $N$  rows, use a **LIMIT** clause:

```
SELECT ename, salary FROM emp
ORDER BY salary DESC LIMIT 5;
```

To start from a row besides the first, use **OFFSET**:

```
SELECT ename, salary FROM emp
ORDER BY salary DESC LIMIT 5 OFFSET 2;
```

- ▶ **LIMIT** improves performance by reducing amount of data processed and sent by the database system.
- ▶ **OFFSET 0** is first row, so **OFFSET 2** would return the 3rd row.

- ▶ **LIMIT/OFFSET** syntax is supported differently by systems.
- ▶ For example, Access uses

```
SELECT TOP 5 eno, salary FROM emp
```

Program	syntax
MySQL, PostgreSQL	LIMIT syntax
Oracle	ROWNUM field that can be filtered in WHERE
SQL Server	SELECT TOP N

(click the pink text above and [here](#) for more details.)

## Try It: SQL `SELECT` with Joins and Ordering

### Example 5.9

Write these three queries:

- ▶ Return all projects with budget < \$500000 sorted by budget descending.
- ▶ List only the top 5 employees by salary descending. Show only their name and salary.
- ▶ List each project pno, dno, pname, and dname ordered by dno ascending then pno ascending. Only show projects if department name > 'D'. Note: This query will require a `join`.



## SELECT Statement Execution Order

Order written:

```
SELECT < feilds >  
FROM < left_table >  
JOIN < right_table >  
ON < join_condition >  
WHERE < where_condition >  
GROUPBY < group_by_list >  
ORDERBY < order_by_list >
```

Order executed:

1. FROM clause
2. JOIN clause
3. WHERE clause
4. GROUP BY clause
5. SELECT clause
6. DISTINCT clause
7. ORDER BY clause

Read more about it [here](#)

## Aggregate Queries and Functions

Several queries cannot be answered using the simple form of the **SELECT** statement. These queries require a summary calculation to be performed. For example:

- ▶ What is the maximum employee salary?
- ▶ What is the total number of hours worked on a project?
- ▶ How many employees are there in department 'D1'?

To answer these queries requires the use of **aggregate** functions. These functions operate on a single column of a table and return a single value.

# Aggregate Functions

Five common aggregate functions are:

**COUNT** returns the number of values in a column

**SUM** returns the sum of the values in a column

**AVG** returns the average of the values in a column

**MIN** returns the smallest value in a column

**MAX** returns the largest value in a column

Notes:

1. **COUNT**, **MAX**, and **MIN** apply to all types of fields, whereas **SUM** and **AVG** apply to only numeric fields.
2. Except for **COUNT(\*)** all functions ignore nulls. **COUNT(\*)** returns the number of rows in the table.
3. Use **DISTINCT** to eliminate duplicates.

## DISTINCT syntax

The following selects the *distinct* salaries from emp:

- ▶ `SELECT DISTINCT salary FROM Emp`

The following counts the number of distinct salaries from emp:

- ▶ `SELECT COUNT(DISTINCT salary) FROM Emp`

WorksOnEdited.odbc : Query1

SALARY
30000.00
50000.00
40000.00
20000.00
50000.00
30000.00
40000.00
50000.00

Record 1 of 8

```
SELECT SALARY FROM EMP;
```

WorksOnEdited.odbc : Query1

SALARY
20000.00
30000.00
40000.00
50000.00

Record 1 of 4

```
SELECT DISTINCT SALARY FROM EMP;
```

WorksOnEdited.odbc : Query1

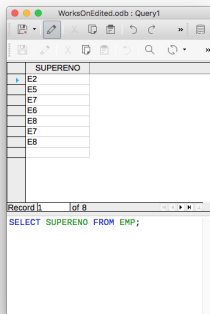
COUNT(EMP.SALARY)
4

Record 1 of 1

```
SELECT COUNT(DISTINCT SALARY)
FROM EMP;
```

## COUNT

Use `SELECT COUNT(*) FROM EMP;` to count all the rows including NULLS and duplicates.

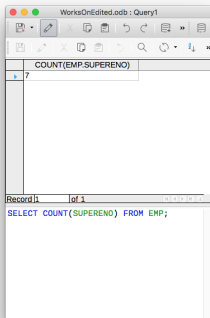


WorksOnEdited.odt : Query1

SUPERENO
E2
E5
E7
E8
E8
E7
E8

Record 1 of 8

```
SELECT SUPERENO FROM EMP;
```

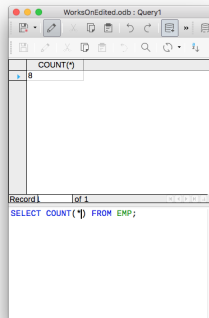


WorksOnEdited.odt : Query1

COUNT(EMP.SUPERENO)
7

Record 1 of 1

```
SELECT COUNT(SUPERENO) FROM EMP;
```



WorksOnEdited.odt : Query1

COUNT(*)
8

Record 1 of 1

```
SELECT COUNT(*) FROM EMP;
```

## SQL

# Aggregate Function Example

Return the number of employees and their average salary.

```
SELECT COUNT(eno) AS numEmp, AVG(salary) AS avgSalary  
FROM emp
```

numEmp	avgSalary
8	38750

Note: **AS** is used to rename a column in the output.

N.B. Aggregate functions are separated by commas just like any other field name.

## GROUP BY Clause

---

Aggregate functions are most useful when combined with the **GROUP BY** clause.

The **GROUP BY** clause groups rows based on the values of the columns specified.

When used in combination with aggregate functions, the result is a table where each row consists of unique values for the **GROUP BY** attributes and the result of the aggregate functions applied to the rows of that group.

## GROUP BY example

For each employee title, return the number of employees with that title.

```
SELECT title, COUNT(eno) AS numEmp  
FROM emp  
GROUP BY title
```

	TITLE	numEmp
▶	EE	2
	SA	3
	ME	2
	PR	1



## GROUP BY example

If a selected field is not aggregated by a function it has to be explicitly added to the GROUP BY clause!!

```
SELECT <A>, >B>, aggregateFun(X) AS varName  
GROUP BY <A>, >B>
```

For example.

```
SELECT title, COUNT(eno) AS numEmp  
FROM emp  
GROUP BY title <- without this line we get an error
```

	TITLE	numEmp
▶	EE	2
	SA	3
	ME	2
	PR	1

## GROUP BY example

For each employee title, return the number of employees with that title, and the minimum, maximum, and average salary.

```
SELECT title, COUNT(eno) AS numEmp,  
       MIN(salary) AS minSal,  
       MAX(salary) AS maxSal, AVG(salary) AS avgSal  
FROM emp  
GROUP BY title
```

title	numEmp	minSal	maxSal	avgSal
EE	2	30000	30000	30000
SA	3	50000	50000	50000
ME	2	40000	40000	40000
PR	1	20000	20000	20000

## GROUP BY Facts

1. You can group by multiple attributes. To be in the same group, all attribute values must be the same.
2. Any **WHERE** conditions are applied before the **GROUP BY** and aggregate functions are calculated.
3. A column name cannot appear in the **SELECT** part of the query unless it is part of an aggregate function or in the list of group by attributes.
4. There is a **HAVING** clause that is applied after the **GROUP BY** clause and aggregate functions are calculated to filter out groups.

## SELECT Statement Execution Order

Order written:

```
SELECT < feilds >  
FROM < left_table >  
JOIN < right_table >  
ON < join_condition >  
WHERE < where_condition >  
GROUPBY < group_by_list >  
HAVING < having_condition >  
ORDERBY < order_by_list >
```

Order executed:

1. FROM clause
2. JOIN clause
3. WHERE clause
4. GROUP BY clause
5. HAVING clause
6. SELECT clause
7. DISTINCT clause
8. ORDER BY clause

Read more about it [here](#)

**WHERE** filters *before* **GROUP BY** whereas **HAVING** filters *after*.

WorksOnitedLab - Query1

</

WorksOnitedLab - Query1

	TITLE	numEmp	minSal	maxSal	avgSal
EE	2	30000	30000	30000	
SA	2	50000	50000	50000	
PR	1	20000	20000	20000	
ME	1	40000	40000	40000	

Record 3 of 4

```
SELECT TITLE,
COUNT(ENO) AS numEmp,
MIN(SALARY) AS minSal,
MAX(SALARY) AS maxSal,
AVG(SALARY) AS avgSal
FROM EMP
WHERE ENAME > 'C'
GROUP BY TITLE
;
```

WorksOnitedLab - Query1

<

## GROUP BY Question

## Example 5.10

Given the emp table and the following query, how many rows are returned?

```
SELECT title, SUM(salary)
FROM emp
GROUP BY title
```

A) 1    B) 2    C) 4    D) 8

emp Table

eno	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

## GROUP BY Question

## Answer

Given the emp table and the following query, how many rows are returned?

```
SELECT title, SUM(salary)
FROM emp
GROUP BY title
```

A) 1    B) 2    C) 4    D) 8

## emp Table

eno	ename	title	salary
E1	J. Doe	EE	30000
E2	M. Smith	SA	50000
E3	A. Lee	ME	40000
E4	J. Miller	PR	20000
E5	B. Casey	SA	50000
E6	L. Chu	EE	30000
E7	R. Davis	ME	40000
E8	J. Jones	SA	50000

TITLE	SUM(EMP.SALARY)
EE	60000
SA	150000
ME	80000
PR	20000

## GROUP BY Question 2

## Example 5.11

Given the workson table and the following query, how many rows are returned?

```
SELECT resp, pno, SUM(hours)
FROM workson
WHERE hours > 10
GROUP BY resp, pno
```

A) 9   B) 7   C) 5   D) 1   E) 0

workson Table

eno	pno	resp	hours
E1	P1	Manager	12
E2	P1	Analyst	24
E2	P2	Analyst	6
E3	P3	Consultant	10
E3	P4	Engineer	48
E4	P2	Programmer	18
E5	P2	Manager	24
E6	P4	Manager	48
E7	P3	Engineer	36

Figure: workson



## GROUP BY Question

## Answer

Given the workson table and the following query, how many rows are returned?

```
SELECT resp, pno, SUM(hours)
FROM workson
WHERE hours > 10
GROUP BY resp, pno
```

A) 9   B) 7   C) 5   D) 1   E) 0

workson Table

eno	pno	resp	hours
E1	P1	Manager	12
E2	P1	Analyst	24
E2	P2	Analyst	6
E3	P3	Consultant	10
E3	P4	Engineer	48
E4	P2	Programmer	18
E5	P2	Manager	24
E6	P4	Manager	48
E7	P3	Engineer	36

RESP	PNO	SUM(WORKSON.HOURS)
Managem	P1	12
Analyst	P1	24
Engineer	P4	48
Program	P2	18
Manager	P2	24
Manager	P4	48
Engineer	P3	36

Record 1 of 7

```
SELECT RESP, PNO, SUM(HOURS)
FROM WORKSON
WHERE HOURS>10
GROUP BY RESP, PNO;
```

## Try It: GROUP BY

### Example 5.12

Use **GROUP BY** and aggregation functions to answer these queries.

1. Output the total number of projects in the database.
2. Return the sum of the budgets for all projects.
3. For each department (dno), return the department number (dno) and the average budget of projects in that department.
4. For each project (pno), return the project number (pno) and the sum of the number of hours employees have worked on that project.
  - ▶ **Challenge:** Show the project name (pname) as well as the project number.
5. **Challenge:** Show the department name (dname), project name (pname), and sum of hours worked on that project as well as the number of employees working on the project.

## Putting it all together

The steps to write an English query in SQL are:

1. Find the columns that you need and put in **SELECT** clause.
2. List the tables that have the columns in the **FROM** clause. If there is more than one, join them together (using **JOIN/ON**)
3. If you must filter rows, add a filter criteria in **WHERE** clause.
4. If you need to create an aggregate, use aggregation functions (e.g. **COUNT**, **AVG**) and **GROUP BY**.
5. If you must filter aggregates, add a filter criteria in a **HAVING** clause.

## Putting it all together

Example: For each project name list the sum of the hours worked by employees working as a 'Manager' on the project.

```
SELECT pname, SUM(hours) as totalHours
FROM workson INNER JOIN proj on workson.pno=proj.pno
WHERE resp='Manager'
GROUP BY pname
```

## Microsoft Access Querying Summary

---

1. Projection is performed by selecting the fields in the output in the field row in the table at the bottom of the screen.
2. Selection is performed by entering the condition in the criteria box. The criteria applies to the field in that column.
3. The tables used are added to the query by the **Show Table...** option.
4. Joins (based on relationships) are often automatically added, but if not, you can add them by selecting the join field in one table, holding the mouse button, then dragging to the join field in the other table.

# Conclusion

---

A **database** is a collection of related data. A database system allows storing and querying a database.

**SQL** is the standard query language for databases, although Microsoft Access also provides a graphical user interface.

**CREATE TABLE** creates a table. **INSERT**, **DELETE**, and **UPDATE** commands modify the data stored within the database.

The basic query operations are selection (subset of rows), projection (subset of columns), join (combine two or more tables), and grouping and aggregation.

# Objectives

- ▶ Define: database, database system, schema, metadata
- ▶ Define: relation, attribute, tuple, domain, degree, cardinality
- ▶ SQL properties: reserved words, case-insensitive, free-format
- ▶ Be able to create a table using CREATE TABLE command and in Microsoft Access.
- ▶ Explain what a key is and what it is used for.
- ▶ Use DROP TABLE to delete a table and its data.
- ▶ Use INSERT/UPDATE/DELETE to add/update/delete rows of a table and perform same actions using Microsoft Access user interface.
- ▶ Execute queries using SQL SELECT and using Microsoft Access user interface.
- ▶ Sort rows using ORDER BY. Use LIMIT to keep only the first (top) N rows.
- ▶ Use GROUP BY and aggregation functions for calculating summary data.
- ▶ Given a small database write simple English queries in SQL.